

MODEL REFERENCE SHEET

Artificial Analysis + CursorBench 3.2 snapshot

26 JUL 2026

13 supplied charts + 50-entry live leaderboard

GENERAL INTELLIGENCE Claude Fable 5 60	AGENTIC GPT-5.6 Sol max 54.0	CODING GPT-5.6 Sol xhigh 78.3
OUTPUT SPEED GPT-5.6 Luna max 185 tok/s	FIRST TOKEN Luna NR / Terra NR 0.7 s	500-TOKEN E2E GPT-5.6 Luna NR 3.5 s
LOWEST COST GPT-5.6 Luna low \$0.04/task	EARLY WRITING Claude Fable 5 max ~2,810 Elo	CURSORBENCH 3.2 Claude Fable 5 max 70.5%

Quick routing guide		
Use case	Primary pick	Why
Maximum overall quality / writing	Claude Fable 5 max	CursorBench leader (70.5%); also leads AA intelligence (60) and early writing (~2,810 Elo). Opus 5 max is 0.5 points behind at 48% of the cost, so reserve Fable max for writing-weighted work.
Best premium quality per dollar	Claude Opus 5 xhigh	69.3% at \$7.35 beats Fable 5 xhigh (68.4%, \$11.73) on both score and cost, and sits 1.2 points under Fable max.
Default quality workhorse	Claude Opus 5 high	66.7% at \$3.91 using 27,932 tokens and 48 steps. Beats GPT-5.6 Sol xhigh by 2.2 points at the same cost.
Maximum non-Anthropic quality	GPT-5.6 Sol max	AA Agentic leader (54.0) and best GPT on CursorBench (67.2%). Use when provider diversity matters or an Anthropic route is unavailable.
Best daily coding quality	GPT-5.6 Sol xhigh	AA Coding leader (78.3); CursorBench 64.5%. Prefer Opus 5 high when ambiguous multi-file repository work dominates.
Fast coding value	GPT-5.6 Sol medium	76.3 AA coding at \$0.31 and 14.3 s E2E; CursorBench 60.0% shows a real quality trade-off versus max, not near-parity.
High-throughput reasoning	GPT-5.6 Luna max	185 tok/s and 51 AA intelligence; CursorBench 61.1% but 87,973 tokens and 61 steps per task.
Interactive balanced default	GPT-5.6 Terra high	49 AA intelligence, 119 tok/s and 6.1 s E2E; CursorBench 54.2% at \$0.89/task.
Lowest-cost strong reasoning	GPT-5.6 Luna xhigh	49 AA intelligence at \$0.14/task; CursorBench 57.7% at \$1.14/task and 48 steps.
Fast simple-task mode	GPT-5.6 Luna non-reasoning	0.7 s TTFT and 3.5 s AA E2E. CursorBench does not list a non-reasoning Luna route; use only for bounded work.
Fast parent-directed delegation	xAI Grok 4.5 high	66.7% at \$1.51/task on CursorBench; 88 tok/s and 16.3 s AA E2E. Cheapest route at this score, but discount the exact rank, not the speed/value signal.

NR = non-reasoning. "Best value" and routing labels are derived judgments from the supplied metrics, not source-chart categories. Values marked ~ elsewhere are estimates from cropped or unlabeled chart elements. Claude Opus 5 and Gemini 3.6 Flash rest on CursorBench evidence only; neither appears in the supplied Artificial Analysis charts.

26 July 2026 refresh: Claude Opus 5 and Gemini 3.6 Flash

The CursorBench 3.2 leaderboard grew from 42 to 50 entries. No previously recorded value changed, and Artificial Analysis has published nothing for either new family, so every claim on this page is CursorBench-only evidence.

			New CursorBench 3.2 entries		
#	Model / effort	Score	Cost / task	Tokens / task	Steps / task
2	Opus 5 Max	70.0%	\$8.23	61,838	78
3	Opus 5 Extra High	69.3%	\$7.35	54,239	72
6	Opus 5 High	66.7%	\$3.91	27,932	48
13	Opus 5 Medium	64.3%	\$3.29	23,612	44
16	Opus 5 Low	62.8%	\$2.55	18,529	37
36	Gemini 3.6 Flash High	53.5%	\$1.56	30,436	64
41	Gemini 3.6 Flash Medium	51.2%	\$1.48	28,511	62
47	Gemini 3.6 Flash Low	47.4%	\$1.13	20,529	50
			What the new entries change		
Finding		Evidence	Action		
Opus 5 resets the premium tier		Opus 5 max 70.0% / \$8.23 against Fable 5 max 70.5% / \$17.32. Opus 5 xhigh 69.3% / \$7.35 beats Fable 5 xhigh 68.4% / \$11.73 on both axes.	Make Opus 5 xhigh the default premium route and keep Fable max for writing-weighted work.		
Opus 5 high is the best quality per dollar		66.7% at \$3.91 using 27,932 tokens in 48 steps, against Sol xhigh at 64.5% and \$3.88.	Prefer Opus 5 high anywhere Sol xhigh was chosen purely for quality.		
Opus 5 high matches Grok 4.5 high without the caveat		Both score 66.7%. Grok costs \$1.51 but carries Cursor's disclosed training-data advantage.	Keep Grok for speed and cost; use Opus 5 high when the evidence must be uncontaminated.		
Opus 4.8 is fully superseded		Opus 5 low 62.8% / \$2.55 beats Opus 4.8 max 62.3% / \$5.77, and every other 4.8 effort loses to a cheaper Opus 5 effort.	Remove Opus 4.8 from routing tables entirely.		
Sonnet 5 is superseded for repository work		Sonnet 5 max 61.5% / \$6.45 with 92,882 tokens and 86 steps against Opus 5 low 62.8% / \$2.55 with 18,529 tokens and 37 steps.	Stop routing ambiguous repository tasks to Sonnet 5 on cost grounds.		
Gemini 3.6 Flash is not competitive at its price		Flash high 53.5% / \$1.56 against Composer 2.5 at 56.1% / \$0.44 and Grok 4.5 low at 63.5% / \$1.22.	No routing change; both cheaper alternatives score higher.		

Opus 5 and Gemini 3.6 Flash have no Artificial Analysis intelligence, agentic, coding, latency, or writing-Elo rows, so their speed and general-capability profiles are unverified in this sheet. CursorBench results are subject to variance and small differences may not be statistically meaningful.

Full visible index rankings

Compact names: Sol/Terra/Luna imply GPT-5.6; "5.5" implies GPT-5.5; NR = non-reasoning. Ties retain the left-to-right order shown in the screenshots.

#	Agentic model	Score	#	Coding model	Score	#	Intelligence model	Score
1	Sol max	54.0	1	Sol xhigh	78.3	1	Fable 5 fallback	60
2	Fable 5 fallback	52.8	2	Sol max	77.4	2	Sol max	59
3	Sol xhigh	51.8	3	Sol high	77.2	3	Sol xhigh	58
4	Sol high	48.5	4	Terra max	76.7	4	Sol high	56
5	Terra max	47.4	5	Fable 5 fallback	76.5	5	Opus 4.8 max	56
6	Opus 4.8 max	47.2	6	Sol medium	76.3	6	Terra max	55
7	Sonnet 5 max	46.7	7	5.5 xhigh	74.9	7	5.5 xhigh	55
8	Grok 4.5 high	45.7	8	Opus 4.8 max	74.3	8	Grok 4.5 high	54
9	Luna max	45.6	9	Grok 4.5 high	72.4	9	Sol medium	54
10	5.5 xhigh	44.9	10	5.5 high	71.6	10	Sonnet 5 max	53
11	Terra xhigh	44.7	11	Sonnet 5 max	71.5	11	5.5 high	53
12	Sol medium	44.5	12	5.5 medium	71.5	12	Terra xhigh	52
13	5.5 high	43.5	13	Luna max	71.4	13	Luna max	51
14	Luna xhigh	42.9	14	Terra xhigh	70.6	14	5.5 medium	50
15	Terra high	41.3	15	Sol low	69.7	15	Sol low	49
16	Luna high	40.1	16	Luna xhigh	68.6	16	Luna xhigh	49
17	Sol low	40.0	17	Terra high	67.1	17	Terra high	49
18	5.5 medium	37.8	18	Sonnet 5 NR	66.4	18	Luna high	46
19	Terra medium	37.0	19	Sol NR	65.1	19	Terra medium	46
20	Sol NR	34.9	20	Terra medium	64.7	20	5.5 low	43
21	Sonnet 5 NR	33.7	21	Luna high	63.3	21	Sonnet 5 NR	42
22	Luna medium	31.0	22	5.5 low	60.9	22	Sol NR	41
23	Terra low	30.6	23	Terra low	58.1	23	Terra low	40
24	5.5 low	30.4	24	5.5 NR	56.5	24	Luna medium	38
25	Terra NR	29.3	25	Terra NR	52.3	25	5.5 NR	35
26	5.5 NR	25.8	26	Luna medium	50.7	26	Terra NR	34
27	Luna NR	~22.0	27	Luna NR	~39.0	27	Luna NR	27

The screenshots show 27 ranked entries in each index. The rightmost Luna non-reasoning Agentic and Coding labels were cropped; they are recorded as approximately 22 and 39.

Model matrix - GPT-5.6 Luna and Terra

Luna maximizes decode throughput; Terra generally provides the stronger interactive quality/latency balance.

Model / mode	Intel	Agent	Code	Cost/task	tok/s	TTFT s	E2E 500 s	Reference note
GPT-5.6 Luna (Non-reasoning)	27	~22.0	~39.0	\$0.05	178	0.7	3.5	Fastest E2E; lowest quality tier.
GPT-5.6 Luna (low)	33	—	—	\$0.04	—	—	—	Cheapest 2D frontier point; perf not shown.
GPT-5.6 Luna (medium)	38	31.0	50.7	\$0.05	177	1.7	4.5	Very low cost; 177 tok/s; 4.5 s E2E.
GPT-5.6 Luna (high)	46	40.1	63.3	\$0.09	176	6.6	9.5	Good low-cost quality step-up.
GPT-5.6 Luna (xhigh)	49	42.9	68.6	\$0.14	174	29.7	32.6	49 intel at \$0.14; long thinking delay.
GPT-5.6 Luna (max)	51	45.6	71.4	\$0.21	185	79.5	82.2	Fastest decode; speed-quality quadrant.
GPT-5.6 Terra (Non-reasoning)	34	29.3	52.3	—	123	0.7	4.8	0.7 s TTFT; cost not visible.
GPT-5.6 Terra (low)	40	30.6	58.1	\$0.10	115	1.3	5.7	Cheap, interactive baseline.
GPT-5.6 Terra (medium)	46	37.0	64.7	\$0.13	120	1.3	5.5	46 intel with 5.5 s E2E.
GPT-5.6 Terra (high)	49	41.3	67.1	\$0.24	119	1.8	6.1	Interactive sweet spot at 49 intel.
GPT-5.6 Terra (xhigh)	52	44.7	70.6	\$0.33	121	16.9	21.1	52 intel; speed-quality quadrant.
GPT-5.6 Terra (max)	55	47.4	76.7	\$0.55	136	133.9	137.6	Strong indexes; 134 s thinking delay.

~ = visually estimated from a clipped chart label. — = the metric was not visible in the supplied screenshots. Cost/task is the Artificial Analysis weighted Intelligence Index task cost, not a token-list price.

Within-family step-up economics			
Upgrade	Quality gain	Cost / time multiplier	Interpretation
Luna xhigh -> max	+2 intel; +2.7 agent; +2.8 code	1.50x cost; 2.52x E2E	Max mainly buys a small quality lift; xhigh is the stronger cost point.
Terra xhigh -> max	+3 intel; +2.7 agent; +6.1 code	1.67x cost; 6.52x E2E	Max meaningfully helps coding, but thinking time explodes from 16.9 s to 133.9 s.
Terra NR -> medium	+12 intel; +7.7 agent; +12.4 code	E2E only 4.8 -> 5.5 s	Medium is a compelling interactive reasoning step-up when cost permits.

Model matrix - GPT-5.6 Sol and GPT-5.5

Sol dominates the top agentic/coding ranks. Medium is the practical coding default; max is an agentic-quality mode with extreme thinking latency.

Model / mode	Intel	Agent	Code	Cost/task	tok/s	TTFT s	E2E 500 s	Reference note
GPT-5.6 Sol (Non-reasoning)	41	34.9	65.1	\$0.20	56	1.0	9.9	Simple tasks; same cost as Sol low.
GPT-5.6 Sol (low)	49	40.0	69.7	\$0.20	51	3.3	13.0	+8 intel vs NR for same \$0.20 cost.
GPT-5.6 Sol (medium)	54	44.5	76.3	\$0.31	53	4.7	14.3	Coding/value sweet spot.
GPT-5.6 Sol (high)	56	48.5	77.2	\$0.45	~49	12.1	22.3	High quality; moderate E2E.
GPT-5.6 Sol (xhigh)	58	51.8	78.3	\$0.68	53	41.6	51.1	Highest Coding Index.
GPT-5.6 Sol (max)	59	54.0	77.4	\$1.04	54	129.3	138.6	Highest Agentic Index.
GPT-5.5 (Non-reasoning)	35	25.8	56.5	\$0.17	66	0.9	8.4	Low latency; low index scores.
GPT-5.5 (low)	43	30.4	60.9	\$0.21	66	1.8	9.4	Small cost/latency step from NR.
GPT-5.5 (medium)	50	37.8	71.5	\$0.41	68	7.4	14.8	50 intel; eclipsed by Sol medium on value.
GPT-5.5 (high)	53	43.5	71.6	\$0.67	70	22.3	29.4	53 intel; 29.4 s E2E.
GPT-5.5 (xhigh)	55	44.9	74.9	\$0.99	68	61.6	69.0	55 intel; \$0.99; 69 s E2E.

~ = visually estimated from a clipped chart label. — = the metric was not visible in the supplied screenshots. Cost/task is the Artificial Analysis weighted Intelligence Index task cost, not a token-list price.

			Sol scale-up: coding saturates before agentic capability		
Mode	Intel	Agentic	Coding	Cost / E2E	Decision cue
low	49	40.0	69.7	\$0.20 / 13.0 s	Same \$0.20 cost as NR, but +8 intelligence.
medium	54	44.5	76.3	\$0.31 / 14.3 s	76.3 coding at \$0.31 and 14.3 s: best coding value.
high	56	48.5	77.2	\$0.45 / 22.3 s	Adds general quality with tolerable 22.3 s E2E.
xhigh	58	51.8	78.3	\$0.68 / 51.1 s	Absolute coding leader; 51.1 s E2E.
max	59	54.0	77.4	\$1.04 / 138.6 s	Agentic leader; only +1.1 coding vs medium at 3.35x cost.

Model matrix - Anthropic and xAI

The Sonnet effort variants are fully visible for speed/latency, but only max and non-reasoning are visible in the index/cost screenshots. Claude Opus 5 does not appear in the supplied Artificial Analysis charts at all, so this sheet carries no AA intelligence, agentic, coding, latency, or speed numbers for it.

Model / mode	Intel	Agent	Code	Cost/task	tok/s	TTFT s	E2E 500 s	Reference note
Claude Sonnet 5 (Non-reasoning)	42	33.7	66.4	\$0.37	63	1.3	9.2	Fast; much lower scores than max.
Claude Sonnet 5 (low)	—	—	—	—	60	1.6	9.8	Latency/speed only in supplied views.
Claude Sonnet 5 (medium)	—	—	—	—	64	2.5	10.3	Latency/speed only in supplied views.
Claude Sonnet 5 (high)	—	—	—	—	65	8.5	16.2	Latency/speed only in supplied views.
Claude Sonnet 5 (xhigh)	—	—	—	—	68	22.7	30.0	Latency/speed only in supplied views.
Claude Sonnet 5 (max)	53	46.7	71.5	\$1.53	83	~193.7	~199.7	53 intel; about 200 s E2E.
Claude Opus 4.8 (max)	56	47.2	74.3	\$1.80	57	34.5	43.2	Premium quality; \$1.80/task.
Claude Fable 5 (with fallback)	60	52.8	76.5	~\$2.70	61	116.2	124.5	Top intelligence/writing; highest cost.
Grok 4.5 (high)	54	45.7	72.4	\$0.31	88	10.6	16.3	54 intel at \$0.31; 88 tok/s.


~ = visually estimated from a clipped chart label. — = the metric was not visible in the supplied screenshots. Cost/task is the Artificial Analysis weighted Intelligence Index task cost, not a token-list price.

				Performance and cost leaderboards			
#	Output speed	tok/s	#	E2E 500 tokens	seconds	#	Lowest cost / task
1	Luna max	185	1	Luna NR	3.5	1	Luna low \$0.04
2	Luna NR	178	2	Luna medium	4.5	2	Luna NR \$0.05
3	Luna medium	177	3	Terra NR	4.8	3	Luna medium \$0.05
4	Luna high	176	4	Terra medium	5.5	4	Luna high \$0.09
5	Luna xhigh	174	5	Terra low	5.7	5	Terra low \$0.10
6	Terra max	136	6	Terra high	6.1	6	Terra medium \$0.13

Output speed measures decode throughput after answering begins. It does not remove reasoning delay: Luna max leads at 185 tok/s yet takes 79.5 s to the first answer token.


Cost-intelligence frontier and the third dimension: time

The supplied synthesis chart treats bubble area as approximate weighted task time. The dashed line is the 2D cost-intelligence efficient frontier.



Early writing benchmark: approximate ranking

The source labels this benchmark as “early results.” Elo and API cost points below are visually read from the axes and rounded; they are not final benchmark values.



Supplied early writing benchmark (source image)

Approximate top results

#	Setting	Cost	Elo
1	Claude Fable 5 - max effort	~\$1.26	~2,810
2	Claude Fable 5 - default	~\$0.35	~2,675
3	Claude Opus 4.8 - max effort	~\$0.61	~2,650
4	GPT-5.6 Sol - default	~\$0.17	~2,550
5	GPT-5.6 Sol - extra high	~\$0.31	~2,485
6	GPT-5.6 Terra - extra high	~\$0.10	~2,440
7	Claude Opus 4.8 - default	~\$0.16	~2,410
8	Grok 4.5 - default	~\$0.02	~2,315
9	GPT-5.5 - extra high	~\$0.40	~2,300
10	GPT-5.6 Luna - extra high	~\$0.03	~2,150

Writing-specific takeaways

- Claude Fable 5 leads: max effort is about 2,810 Elo; default is about 2,675.
- Fable max costs about 3.6x default for roughly +135 Elo.
- Claude Opus 4.8 max reaches about 2,650 Elo at about \$0.61/task.
- Grok 4.5 default is unusually inexpensive (~\$0.02) for roughly 2,315 Elo.
- GPT-5.6 Sol default appears to beat its extra-high setting while costing less; treat this non-monotonic result as provisional.
- GPT-5.6 Terra extra high reaches roughly 2,440 Elo at about \$0.10/task.
- Writing effort labels do not necessarily map one-to-one to the index-chart effort labels.
- Writing API cost/task is a different workload from Intelligence Index cost/task; do not compare the dollar values directly.

Metric definitions, cost anatomy, and transcription notes

Use this page when applying the sheet to routing, procurement, or benchmark interpretation.

Definitions Intelligence Index v4.1: Composite of 9 evaluations: GDPval-AA v2, tau^3-Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, and AA-LCR. Agentic Index: Average of agentic-capability benchmarks in the source: GDPval-AA v2 and tau^3-Banking. Coding Index: Weighted average of Terminal-Bench v2.1 and SciCode. Output speed: Output tokens per second after generation begins; higher is better. TTFT: Seconds to first answer token, including reasoning-model thinking time; lower is better. E2E 500: Seconds to produce 500 output tokens, including thinking and input processing; lower is better. Cost/task: Weighted average USD per Artificial Analysis Intelligence Index task; lower is better. Writing Elo: Separate early writing benchmark rating plotted against API cost per writing task. CursorBench 3.2: Correctness score for ambiguous, multi-file tasks sourced from real Cursor sessions. The suite includes instruction-following and advanced tool-use problems; higher is better. CursorBench cost/task: Published input, cache-read, cache-write, and output prices applied to each task's token use, then averaged with the score weights. It is not comparable to AA cost/task because the workloads differ.	
Cost anatomy in the supplied charts - Token-type stack: answer, reasoning, cache write, cache hit, and input. - Evaluation stack: GDPval-AA v2, tau^3-Banking, Terminal-Bench v2.1, SciCode, Humanity's Last Exam, GPQA Diamond, CritPt, AA-Omniscience, and AA-LCR. - At the high-cost end, GDPval-AA v2 is visibly a major contributor; Terminal-Bench is also material for Claude Opus/Fable. - Use the total bar labels as the primary comparable number. Several segment labels are not visible in the screenshot. Interpretation cautions - Index values are aggregate scores, not percentages or guaranteed task success rates. - The UI screenshots are horizontally scrollable; some selected models may be outside the captured viewport. - Latency is environment- and provider-dependent. Treat it as the measured snapshot, not a service-level guarantee. - Fable is "with fallback" in index charts, while the writing chart uses default/max effort labels; these may not be identical routes. - Reasoning modes are indicated by a lightbulb in the source charts. Non-reasoning rows are explicitly labeled NR here. - The weighted task-time bubbles in the frontier chart are not the same measure as the 500-token E2E bars. CursorBench caution: Results are subject to variance and small score differences may not be statistically meaningful. Grok 4.5 has a disclosed training-data advantage on this suite, so its exact rank is discounted; independent AA speed, cost, agentic, and coding results still support a fast/value role.	
Source image coverage	
Supplied image	What was extracted
Artificial Analysis Agentic Index	Full 27-entry agentic ranking and scores.
Artificial Analysis Coding Index	Full 27-entry coding ranking and scores.
Artificial Analysis Intelligence Index	Full 27-entry intelligence ranking and benchmark composition.
Intelligence vs. Output Speed	Speed-quality quadrant and model trade-offs.
End-to-End Response Time	500-token completion times and thinking/output decomposition.
Latency - Time To First Answer Token	First-token latency including reasoning time.
Output Speed	Decode throughput in output tokens per second.
Evaluation Breakdown	Cost contribution by each Intelligence Index evaluation.
Cost per Intelligence Index Task	Total cost and token-type cost stack.
Intelligence vs. Cost per Intelligence Index Task	Cross-vendor cost-quality quadrant.
HM8OUrWWMAAkIO1.jpeg	Early writing Elo versus API task cost.
HM6BgeuaMAA6XR8.jpeg	OpenAI family cost-intelligence curves and attractive quadrant.
bench.jpeg	2D efficient frontier plus weighted task-time bubbles.
cursor.com/cursorbench (live, 26 Jul 2026)	Full 50-entry CursorBench 3.2 leaderboard, including the new Claude Opus 5 and Gemini 3.6 Flash entries.

Transcription rule: exact Artificial Analysis labels were copied from the supplied images; values prefixed with ~ were estimated from chart geometry or cropped labels. CursorBench 3.2 values were refreshed from the live leaderboard on 26 July 2026. Sources: <https://cursor.com/cursorbench> and <https://cursor.com/blog/cursorbench>.

CursorBench 3.2 - full leaderboard (1-25)

Cursor evaluates agents on ambiguous, multi-file tasks from real Cursor sessions. Version 3.2 added instruction-following and advanced tool-use problems. The live table was refreshed 26 July 2026 and now lists 50 entries: Claude Opus 5 (five effort levels) and Gemini 3.6 Flash (three) joined the 42 entries recorded on 18 July, and no previously recorded score, cost, token, or step value changed.

#	Model / effort	Score	Cost / task	Tokens / task	Steps / task
1	Fable 5 Max	70.5%	\$17.32	103,525	72
2	Opus 5 Max	70.0%	\$8.23	61,838	78
3	Opus 5 Extra High	69.3%	\$7.35	54,239	72
4	Fable 5 Extra High	68.4%	\$11.73	64,971	56
5	GPT-5.6 Sol Max	67.2%	\$5.69	28,320	48
6	Opus 5 High	66.7%	\$3.91	27,932	48
7	Grok 4.5 High*	66.7%	\$1.51	19,521	33
8	Fable 5 High	66.5%	\$8.77	43,747	48
9	Grok 4.5 Medium*	65.4%	\$1.54	18,914	34
10	Fable 5 Medium	65.2%	\$6.80	30,366	41
11	GPT-5.6 Terra Max	64.9%	\$2.89	32,969	47
12	GPT-5.6 Sol Extra High	64.5%	\$3.88	19,699	38
13	Opus 5 Medium	64.3%	\$3.29	23,612	44
14	Grok 4.5 Low*	63.5%	\$1.22	15,841	31
15	GPT-5.6 Sol High	63.5%	\$2.79	13,867	32
16	Opus 5 Low	62.8%	\$2.55	18,529	37
17	Opus 4.8 Max	62.3%	\$5.77	71,411	44
18	Fable 5 Low	62.1%	\$4.46	18,182	31
19	Sonnet 5 Max	61.5%	\$6.45	92,882	86
20	GPT-5.6 Luna Max	61.1%	\$1.97	87,973	61
21	GPT-5.6 Sol Medium	60.0%	\$1.95	9,747	27
22	Opus 4.8 Extra High	59.4%	\$4.50	51,121	40
23	GPT-5.6 Terra Extra High	59.2%	\$1.44	16,089	29
24	Sonnet 5 Extra High	58.7%	\$4.16	52,871	67
25	GPT-5.5 High	58.4%	\$2.05	12,183	28

Top-line result: Fable 5 max still leads at 70.5%, but Claude Opus 5 max is now second at 70.0% for \$8.23 against Fable max's \$17.32. Opus 5 extra high is third at 69.3% and \$7.35, ahead of Fable 5 extra high on both score and cost.

CursorBench 3.2 - full leaderboard (26-50)

#	Model / effort	Score	Cost / task	Tokens / task	Steps / task
26	GPT-5.5 Extra High	58.4%	\$2.85	17,534	32
27	Opus 4.8 High	58.0%	\$3.15	33,548	33
28	GPT-5.6 Luna Extra High	57.7%	\$1.14	22,480	48
29	Sonnet 5 High	56.9%	\$3.19	39,483	57
30	GPT-5.6 Luna High	56.8%	\$0.82	15,141	40
31	Opus 4.8 Medium	56.1%	\$2.81	28,384	32
32	Composer 2.5	56.1%	\$0.44	14,286	33
33	GLM 5.2 Max	55.0%	\$1.76	35,946	58
34	GPT-5.6 Terra High	54.2%	\$0.89	9,468	23
35	GPT-5.5 Medium	53.8%	\$1.51	8,522	25
36	Gemini 3.6 Flash High	53.5%	\$1.56	30,436	64
37	Opus 4.8 Low	53.1%	\$2.02	19,624	27
38	GPT-5.6 Sol Low	52.6%	\$1.01	5,104	19
39	Sonnet 5 Medium	52.4%	\$2.16	26,200	46
40	GLM 5.2 High	51.5%	\$1.19	21,829	49
41	Gemini 3.6 Flash Medium	51.2%	\$1.48	28,511	62
42	GPT-5.6 Terra Medium	50.3%	\$0.61	6,222	20
43	Kimi K2.7 Code	49.7%	\$1.43	31,247	58
44	Gemini 3.5 Flash	48.8%	\$2.20	46,702	77
45	GPT-5.6 Luna Medium	47.7%	\$0.39	7,095	28
46	Sonnet 5 Low	47.7%	\$1.30	16,269	33
47	Gemini 3.6 Flash Low	47.4%	\$1.13	20,529	50
48	GPT-5.6 Terra Low	46.9%	\$0.53	5,312	19
49	GPT-5.5 Low	46.6%	\$0.98	5,168	20
50	GPT-5.6 Luna Low	37.6%	\$0.16	3,209	17

* Cursor discloses that Grok 4.5 benefited because an earlier Cursor codebase snapshot was unintentionally present in training. The impact is unknown, so treat its exact 66.7% score cautiously. The separate AA results still establish its speed, low cost, and competitive general capability.

Cost/task uses published token prices and this benchmark's task weights. Results are subject to variance; small score differences may not be statistically meaningful. Compare models within CursorBench 3.2, not against percentages from another suite.

Cross-benchmark synthesis and revised routing

Artificial Analysis and CursorBench answer different questions. AA separates general intelligence, agentic capability, coding, latency, decode speed, and cost. CursorBench is narrower but more directly aligned with ambiguous repository work. No combined score is calculated because the scales and workloads are not commensurate.

Model	AA Intel / Agent / Code	CursorBench score / cost	Interpretation
Fable 5 max / fallback†	60 / 52.8 / 76.5	70.5% / \$17.32	Still the CursorBench and AA general-quality leader, but only 0.5 points above Opus 5 max at 2.1x the cost.
Claude Opus 5 max	Not in supplied AA charts	70.0% / \$8.23	New premium value leader: within CursorBench variance of Fable max at 48% of its cost.
Claude Opus 5 xhigh	Not in supplied AA charts	69.3% / \$7.35	Strictly better than Fable 5 xhigh (68.4% / \$11.73) on both score and cost.
Fable 5 extra high	Not separately visible	68.4% / \$11.73	Superseded by Opus 5 xhigh, which scores higher for \$4.38 less per task.
GPT-5.6 Sol max	59 / 54.0 / 77.4	67.2% / \$5.69	Best GPT and AA agentic leader; now the strongest non-Anthropic quality route.
Claude Opus 5 high	Not in supplied AA charts	66.7% / \$3.91	Best quality per dollar: +2.2 points over Sol xhigh at the same cost, in 48 steps.
Grok 4.5 high*	54 / 45.7 / 72.4	66.7% / \$1.51	Cheapest route at this score; exact rank discounted, AA speed and value evidence retained.
GPT-5.6 Sol xhigh	58 / 51.8 / 78.3	64.5% / \$3.88	Keeps the AA coding lead, but Opus 5 high beats it on CursorBench at effectively equal cost.
GPT-5.6 Sol high	56 / 48.5 / 77.2	63.5% / \$2.79	Best non-Anthropic value under \$3/task; 1.0 point below Sol xhigh.
Claude Opus 4.8 (all efforts)	56 / 47.2 / 74.3	62.3% / \$5.77 at max	Fully superseded: Opus 5 low scores 62.8% at \$2.55, and every 4.8 effort loses to a cheaper Opus 5 effort.
GPT-5.6 Sol medium	54 / 44.5 / 76.3	60.0% / \$1.95	Fast AA coding value, but 10.0 points below Opus 5 max on CursorBench.

† AA reports a Fable 5 with-fallback route, while CursorBench reports explicit effort levels. Treat the cross-suite row as family-level evidence, not a controlled effort comparison. Claude Opus 5 and Gemini 3.6 Flash have no rows in the supplied AA charts, so their entries here are CursorBench-only.

Recommendation changes

Use case	Recommendation	Change and rationale
Absolute maximum quality	Claude Fable 5 max	Unchanged, but narrowly. It keeps CursorBench rank 1 and AA writing leadership while Opus 5 max sits 0.5 points behind at 48% of the cost.
Best premium value	Claude Opus 5 xhigh	New. 69.3% at \$7.35 beats Fable 5 xhigh on both axes and lands 1.2 points under Fable max.
Default quality workhorse	Claude Opus 5 high	New. 66.7% at \$3.91 in 48 steps; +2.2 points over Sol xhigh at the same cost.
Maximum non-Anthropic quality	GPT-5.6 Sol max	Reframed. Still the best GPT route and the one to use when Anthropic capacity or provider diversity is the constraint.
Daily responsive coding	GPT-5.6 Sol xhigh	Qualified. Retains the AA coding lead (78.3), but choose Opus 5 high when ambiguous repository work dominates.
Fast parent-directed delegation	xAI Grok 4.5 high	Unchanged. Still the cheapest 66.7% route for scouting, context building, bounded fixes, and worker implementation.
Independent dual review	Opus 5 xhigh + GPT-5.6 Sol xhigh	Updated. Pair the strongest route from each vendor so the two-reviewer gate keeps genuine model diversity.
Writing and long-form prose	Claude Fable 5	Unchanged. Fable leads the AA early-writing Elo (~2,810 at max effort); CursorBench measures repository tasks, not prose.
Retire from routing	Claude Opus 4.8 and Sonnet 5	New. Every Opus 4.8 and Sonnet 5 effort is beaten by a cheaper Opus 5 effort on CursorBench.

Practical conclusion: Claude Opus 5 is the consequential change in this refresh. Use Opus 5 high as the default quality route, Opus 5 xhigh when a few more points matter, and Fable 5 max only when writing quality or the last 0.5 CursorBench points justify 2.1x the cost. Keep xAI Grok 4.5 high for fast parent-directed scouting, context building, bounded fixes, and worker implementation, and keep GPT-5.6 Sol on the review path so the dual-reviewer gate stays cross-vendor.